



Java Serialization for Big Data

Daniel Hinojosa

About Me..

Daniel Hinojosa
Programmer, Consultant, Trainer

On O'Reilly:

Testing in Scala (Book)

Beginning Scala Programming (Video)

Scala Beyond the Basics (Online Training)

TDD in Java (Online Training)


Getting Starting with Scala* (Online Training)

Scala The Next Level* (Online Training)

Contact:

dhinojosa@evolutionnext.com

@dhinojosa





<https://github.com/dhinojosa/serialization-study>



Rationale For Presentation

- Java Serialization is Slow
- Backwards Compatibility is an issue
- Many modern messaging and Big Data aggregation tools use these built in serialization tools
- Cross-language may also be an issue



Apache Avro



Apache Avro

- Created by Doug Cutting; Creator of Hadoop
- Serialization is defined by schema
- Schemas are JSON Based
- Codegen available at command line
- Codegen available by Maven Plugin
- Supports the following languages
 - C, C++, Java, Perl, Python, Ruby, PHP



Avro Name Origin

- Original British Aircraft Manufacturer WWI, WWII
- A.V. Roe and Company established 1910





Apache Avro Types

- Generic Records
 - Develop code that reflects your schema
- Reflection
 - Auto-create schema from an existing class
- Specific
 - Codegen your class from schema
 - This is the most common form



Avro Primitive Types

Avro Type	Java Type
null	null
double	double
float	float
int	int
long	long
bool	boolean
string	Unicode CharSequence
bytes	Sequence of 8-bit unsigned bytes



Avro Complex Types

```
{
  "namespace": "com.evolutionnext.avro"
  "type": "record",
  "doc" : "An music album",
  "name": "Album",
  "fields": [
    {
      "name": "name",
      "type": "string"
    },
    {
      "name": "yearReleased",
      "type": [
        "int",
        "null"
      ]
    }
  ]
}
```

← Package

← Documentation

← Name

← Fields

← Union Type



Avro Field Options

Avro Field	Description
name	Name of the field
doc	Documentation
type	Type of the field
default	Default value
order	What order does this impact record?
aliases	Other names for the field



Avro Enum

```
{ "type" : "enum",  
  "name" : "rainbowColors",  
  "doc" : "Colors of the Rainbow",  
  "symbols" : [ "RED", "ORANGE", "YELLOW",  
                 "GREEN", "BLUE", "INDIGO",  
                 "VIOLET" ] }
```




Avro Tools

```
java -jar avro-tools-1.8.2.jar compile schema <schema file> <destination>
```



avro-tools jar



your avsc file



destination



Apache Avro Demo



Google Protobuf



About Protobuf

- Developed at Google
- Full Name: Protocol Buffer
- Uses Proprietary Language
 - proto2
 - proto3
- Support for
 - C++
 - Java
 - Python
 - Go
 - Ruby
 - C#



Protobuf Installation

- Download Protocol Buffer
- Ensure that g++ Compiler Installed
- Extract Protocol Buffer
- Run the following:
 - `./configure`
 - `make`
 - `make check`
 - `sudo make install`
 - `protoc --version`

For Windows: <https://github.com/google/protobuf/blob/master/src/README.md#c-installation---windows>



Protobuf Simple Types

Proto Type	Java Type
double	double
float	float
sint32	int
sint64	long
bool	boolean
string	String
bytes	ByteString (protobuf)

Other unsigned types are available:

<https://developers.google.com/protocol-buffers/docs/proto>



Using Protoc

```
protoc -I=$SRC_DIR --java_out=$DST_DIR $SRC_DIR/yourproto.proto
```



Where to find imports



Where to output java code



Your definition



Protobuf Build Tool Plugins


- Maven: <https://www.xolstice.org/protobuf-maven-plugin/>
- Gradle: <https://github.com/google/protobuf-gradle-plugin>
- SBT: <https://github.com/sbt/sbt-protobuf>
- Leiningen: <https://github.com/ninjudd/lein-protobuf>



Google Protocol Buffers Demo




Colfer




Colfer

- <https://github.com/pascaldekloe/colfer>
- Fastest Serialization according to benchmarks
- Per documentation: “Suboptimal performance is treated like a bug.”
- Schema/Codegen Based
- Inspired by Google Protocol Buffers
- Codegen can be performed by command line or build tool plugin



Colfer


- Multiple Language Support
 - C, C++
 - Go
 - Java
 - JavaScript



Colfer Simple Types

Colfer Type	Java Type
double	double
float	float
int32	int
int64	long
bool	boolean
timestamp	java.time.Instant
text	String
binary	byte[]

Other unsigned types are available:
<https://github.com/pascaldekloe/colfer>



Colfer Complex Types

//Comments ← Comments appear in generated code


package datebook ← Package

```
type appointment struct {  
    id sint  
    datetime timestamp  
    duration int32  
    asset asset  
}
```

Complex Type


```
type asset struct {  
    name string  
    type string  
}
```

Reference to complex type



Colfer Compiler

- Using **go** run the following to obtain Colfer:
 - `go get -u github.com/pascaldekloe/colfer/cmd/colf`
- To Run:
 - `colf -p com/xyzcorp Java template_dir`
 - `package` → `com/xyzcorp`
 - `language` → `Java`
 - `files/directory of templates` → `template_dir`



Colfer Maven Plugin

- Due to relative obscurity of Colfer, there is only a maven plugin
- Perhaps for other plugins, configuring to execute the colf command line would be in order



Colfer Demo

Care with CodeGen


Many of the Codegen Polyglot Serializable solutions (Protobuf, have unsigned int, long but Java does not support unsigned values!

Care with Serializable

Most will not offer subclassing (which makes sense) and will likely be up to you to reconstitute your object hierarchy.



Kryo



Kryo

- <https://github.com/EsotericSoftware/kryo>
- Non-Schema, Full-Java Serialization Library
- Easy to Use



Kyro Demo

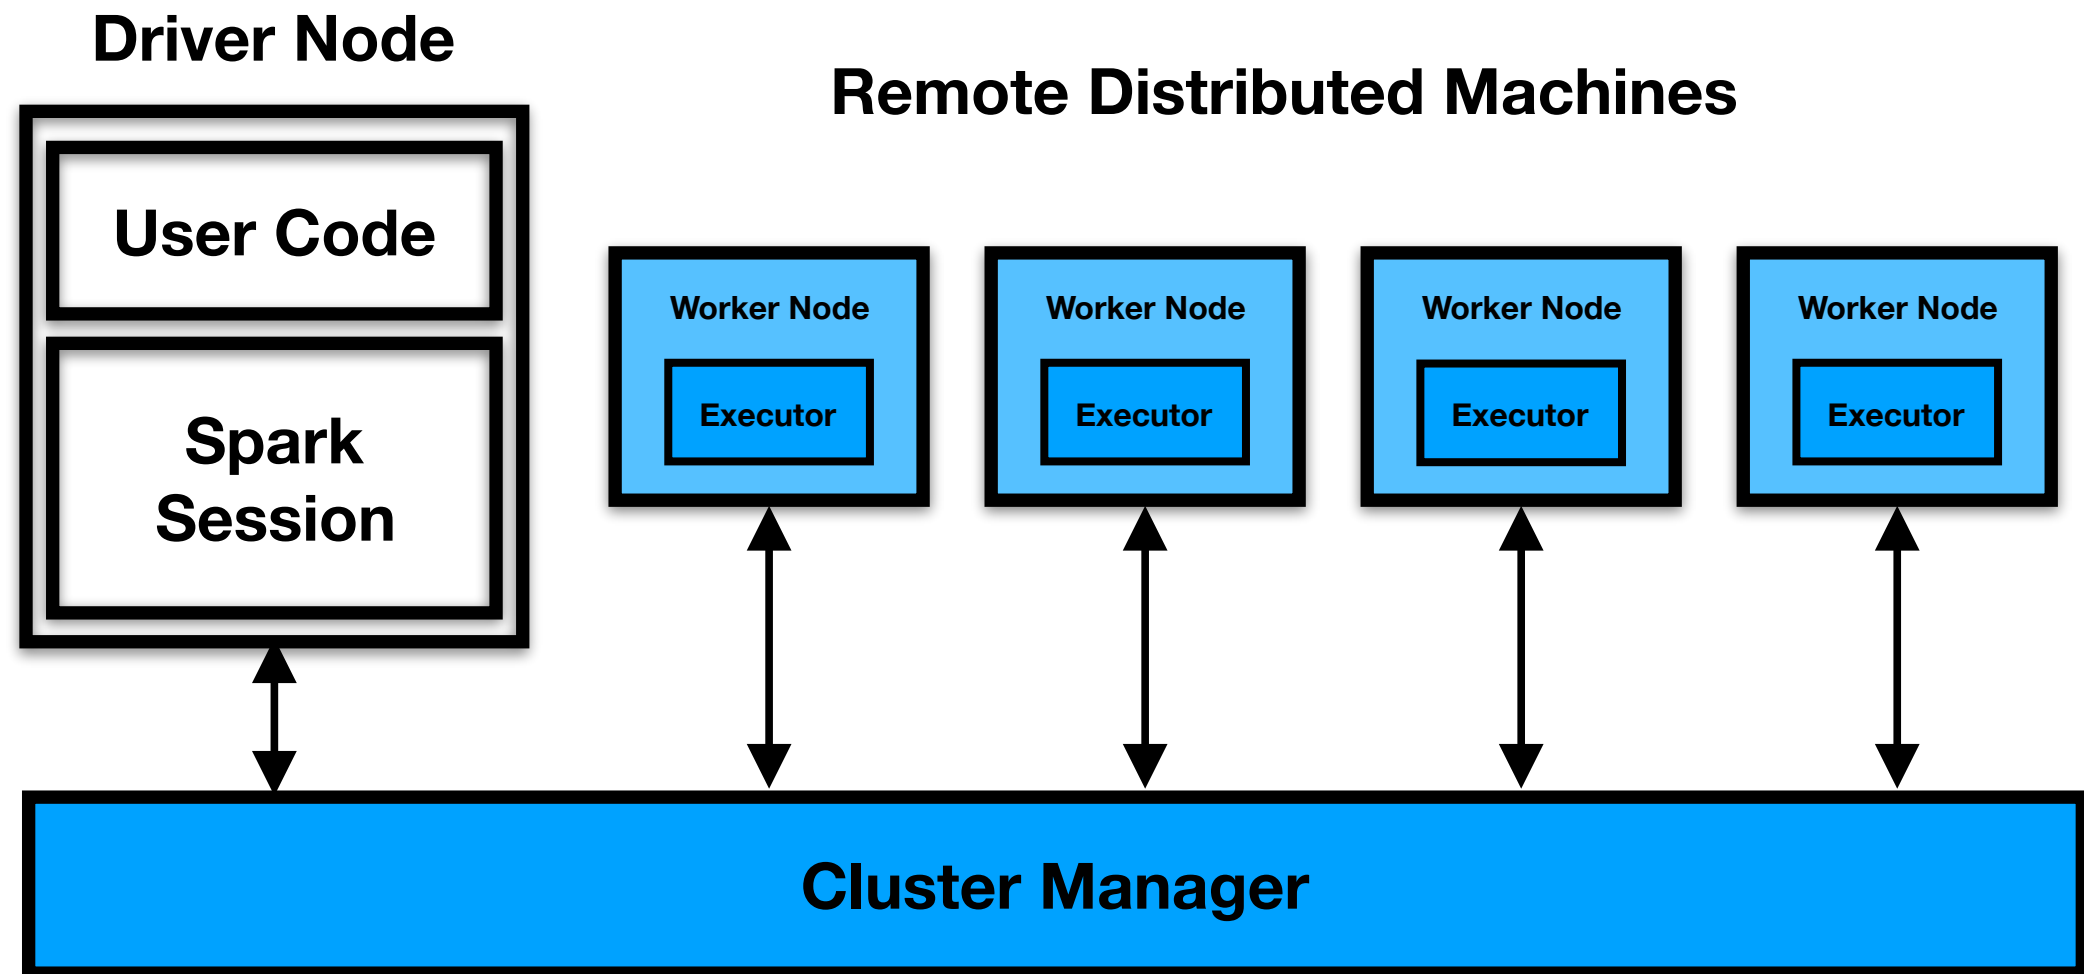


Benchmarks

<https://github.com/eishay/jvm-serializers/wiki>

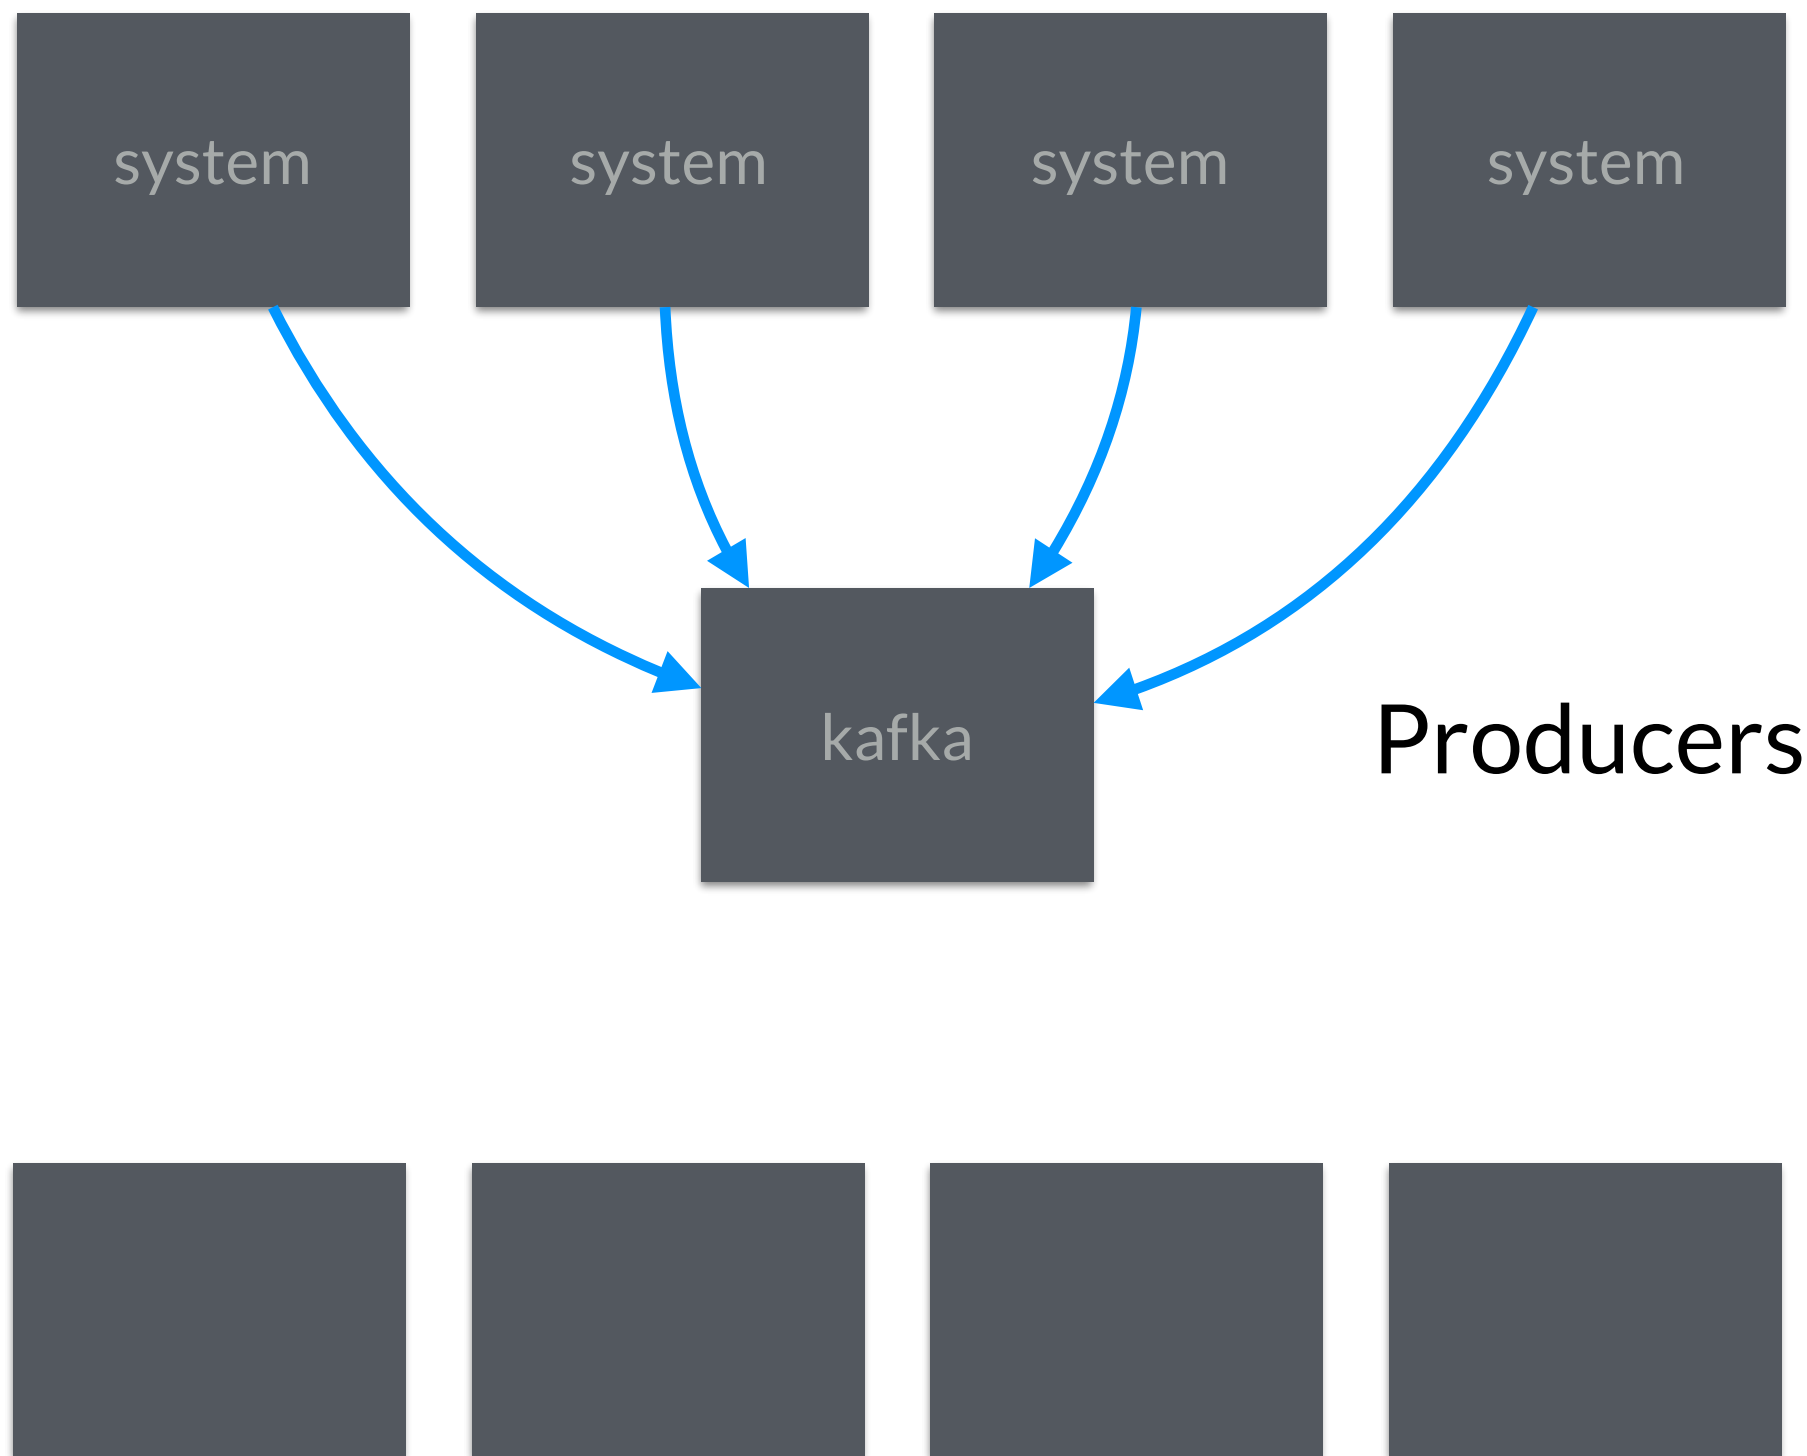




Kryo and Spark




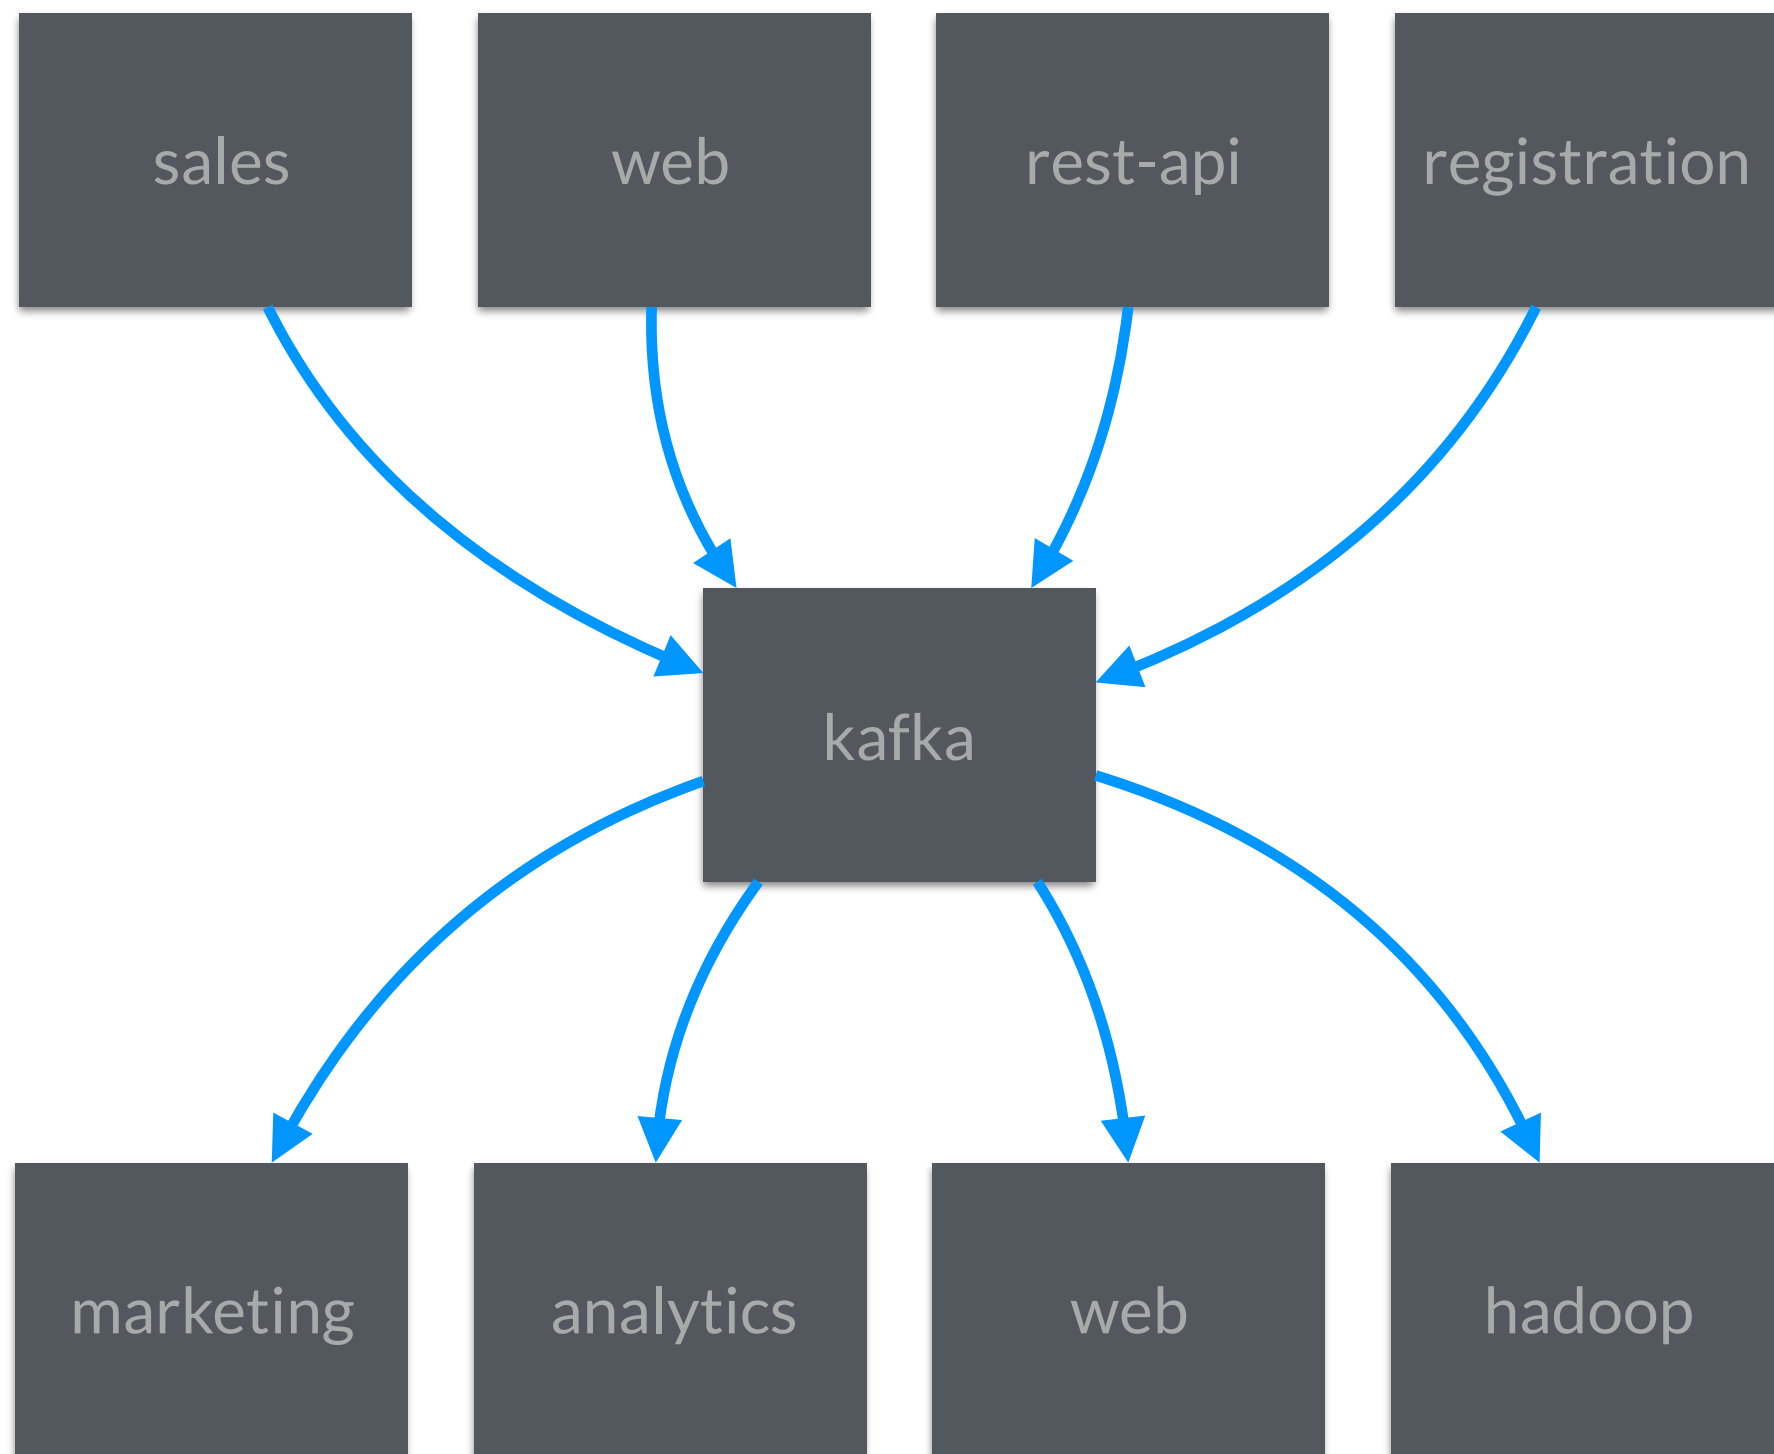


Avro and Kafka













Oversimplified: A Producer can be a Consumer



Kafka Producer

kafka broker: 0

Partition 0:



[0]

[1]


[2]

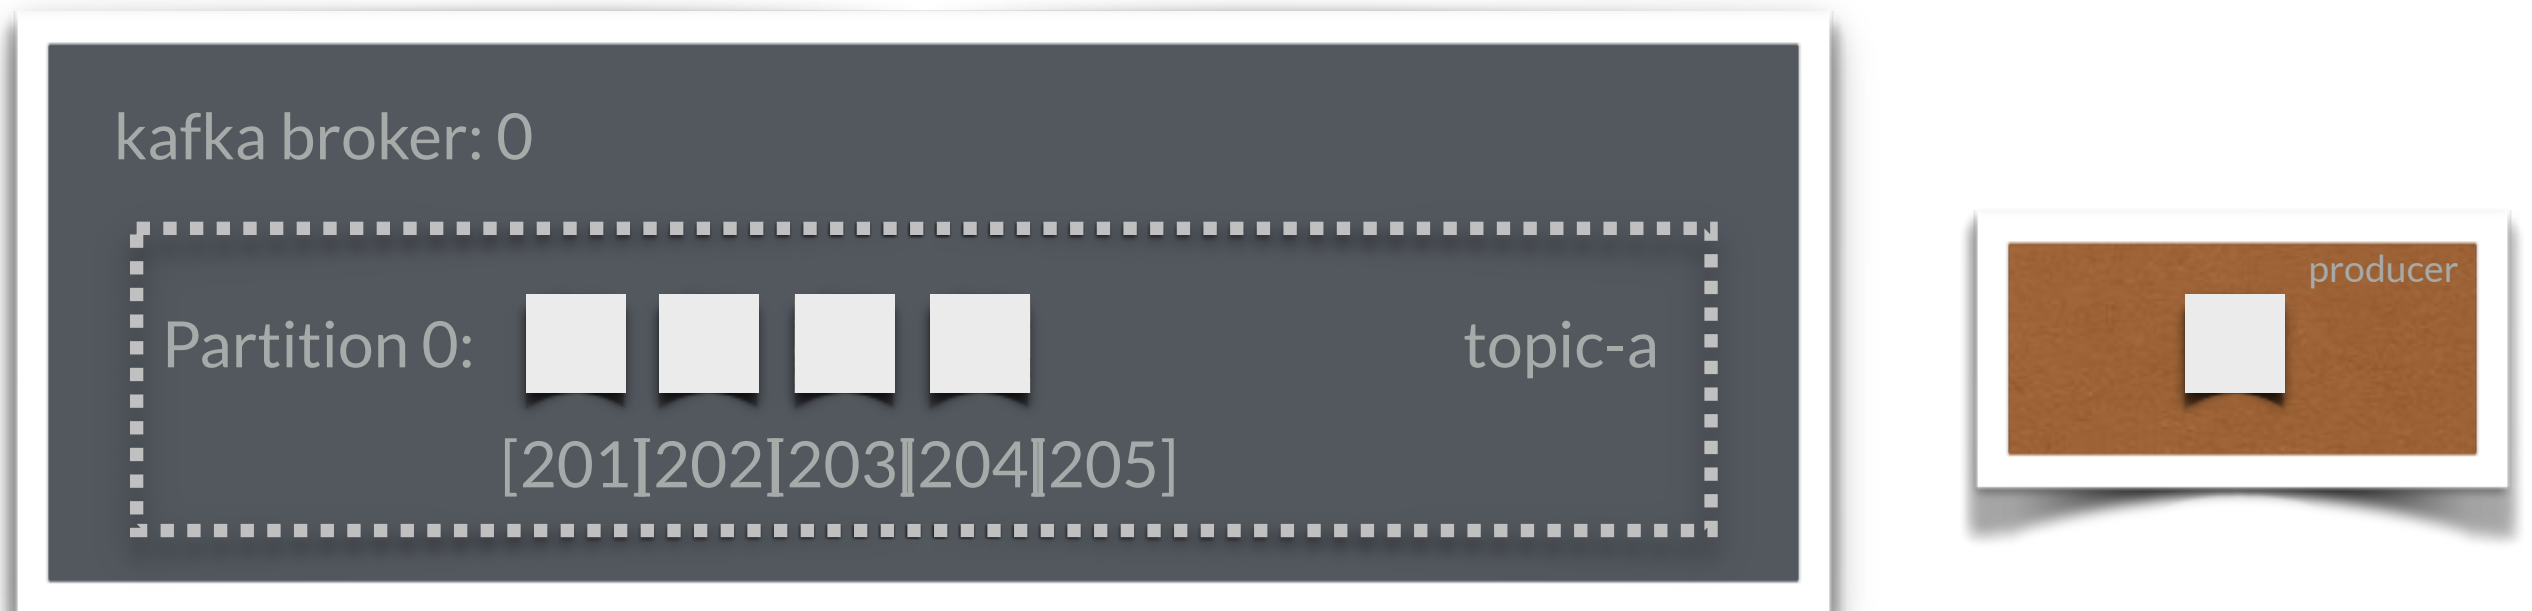
[3]

[4]

topic-a

producer





Retention: The data is temporary




Kafka Consumer


consumer-1_offset

kafka broker: 0


Partition 0:



[0]



[1]



[2]






topic-a

consumer-1

consumer-1_offset

kafka broker: 0

Partition 0:



[0]

[1]

[2]


[3]

[4]

topic-a

consumer-1

from-beginning




from-end

kafka broker: 0


Partition 0:




[0]




[1]



[2]



[3]


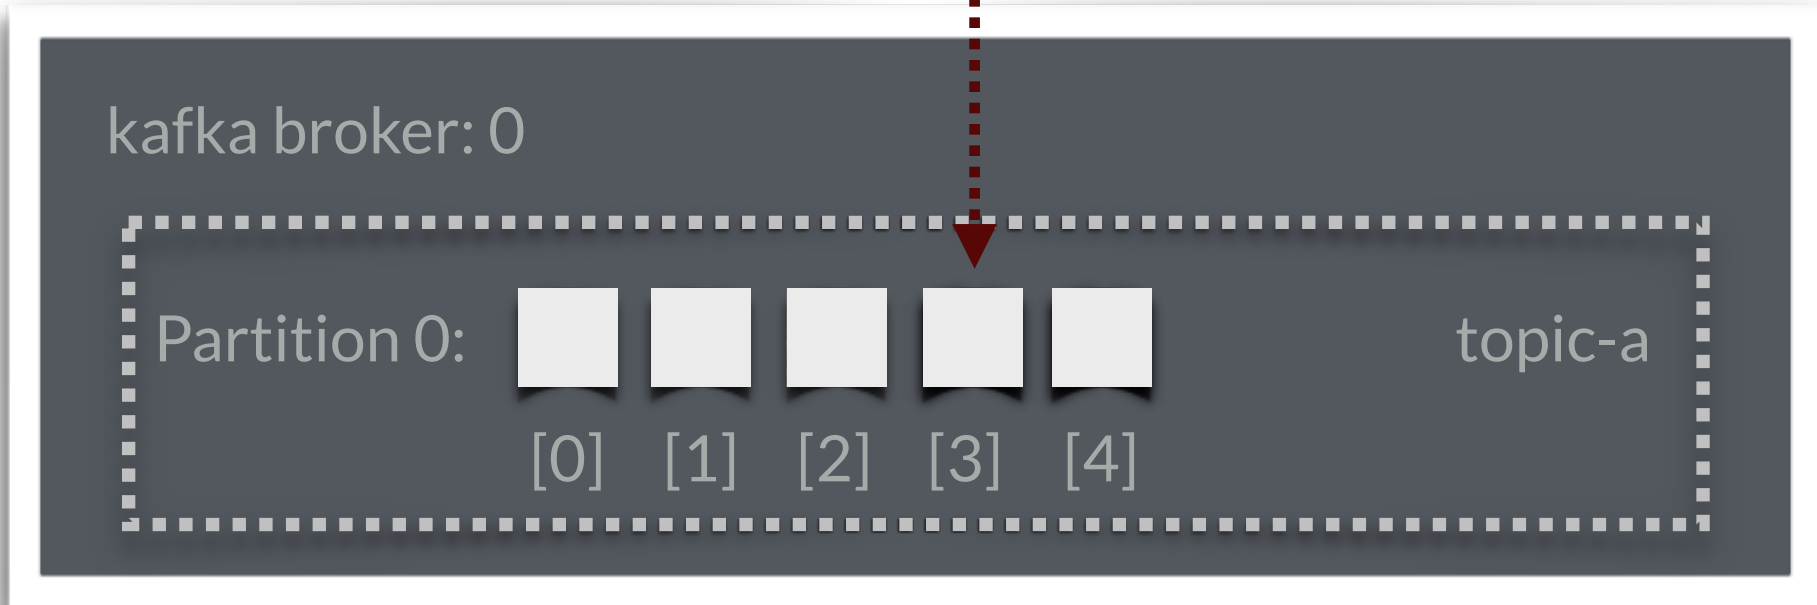


[4]

topic-a

consumer-1

from-offset

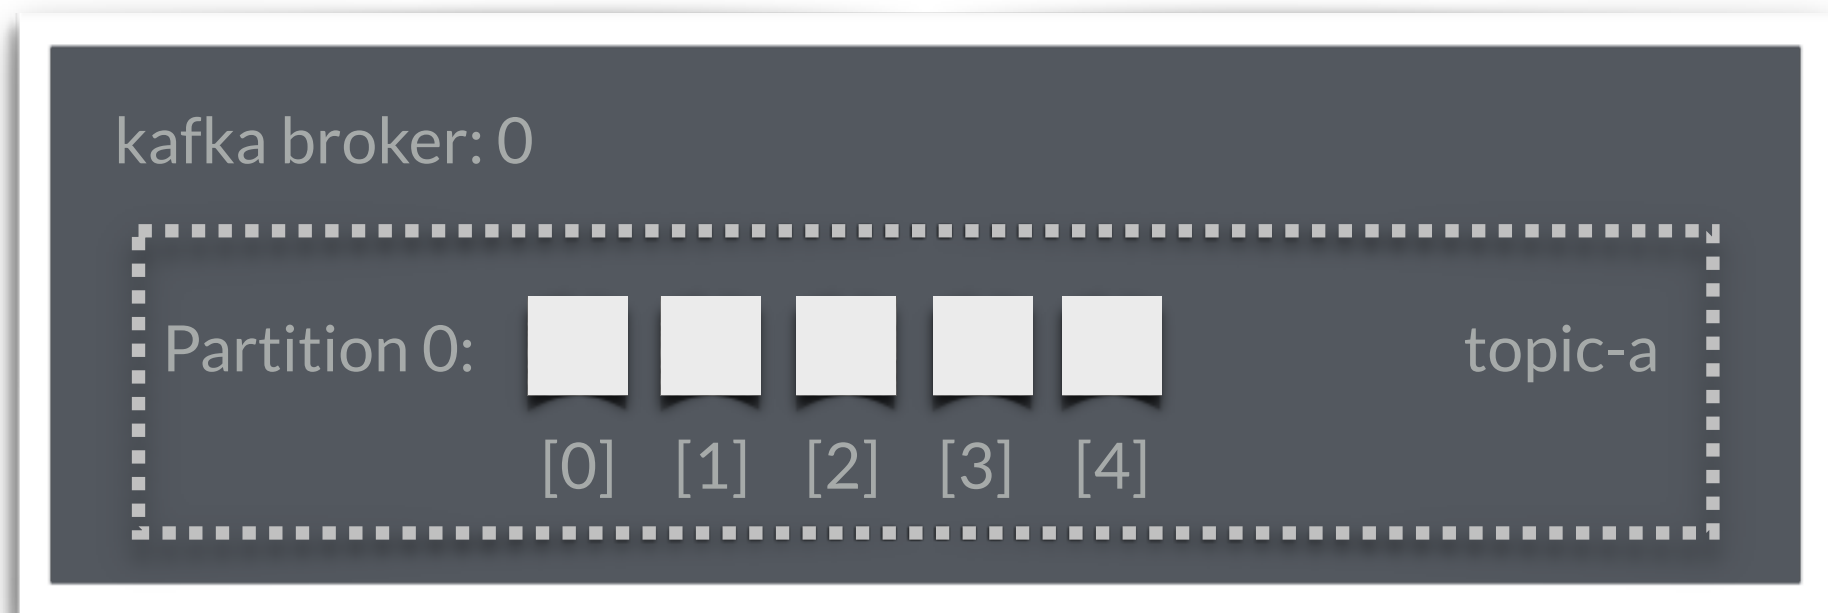
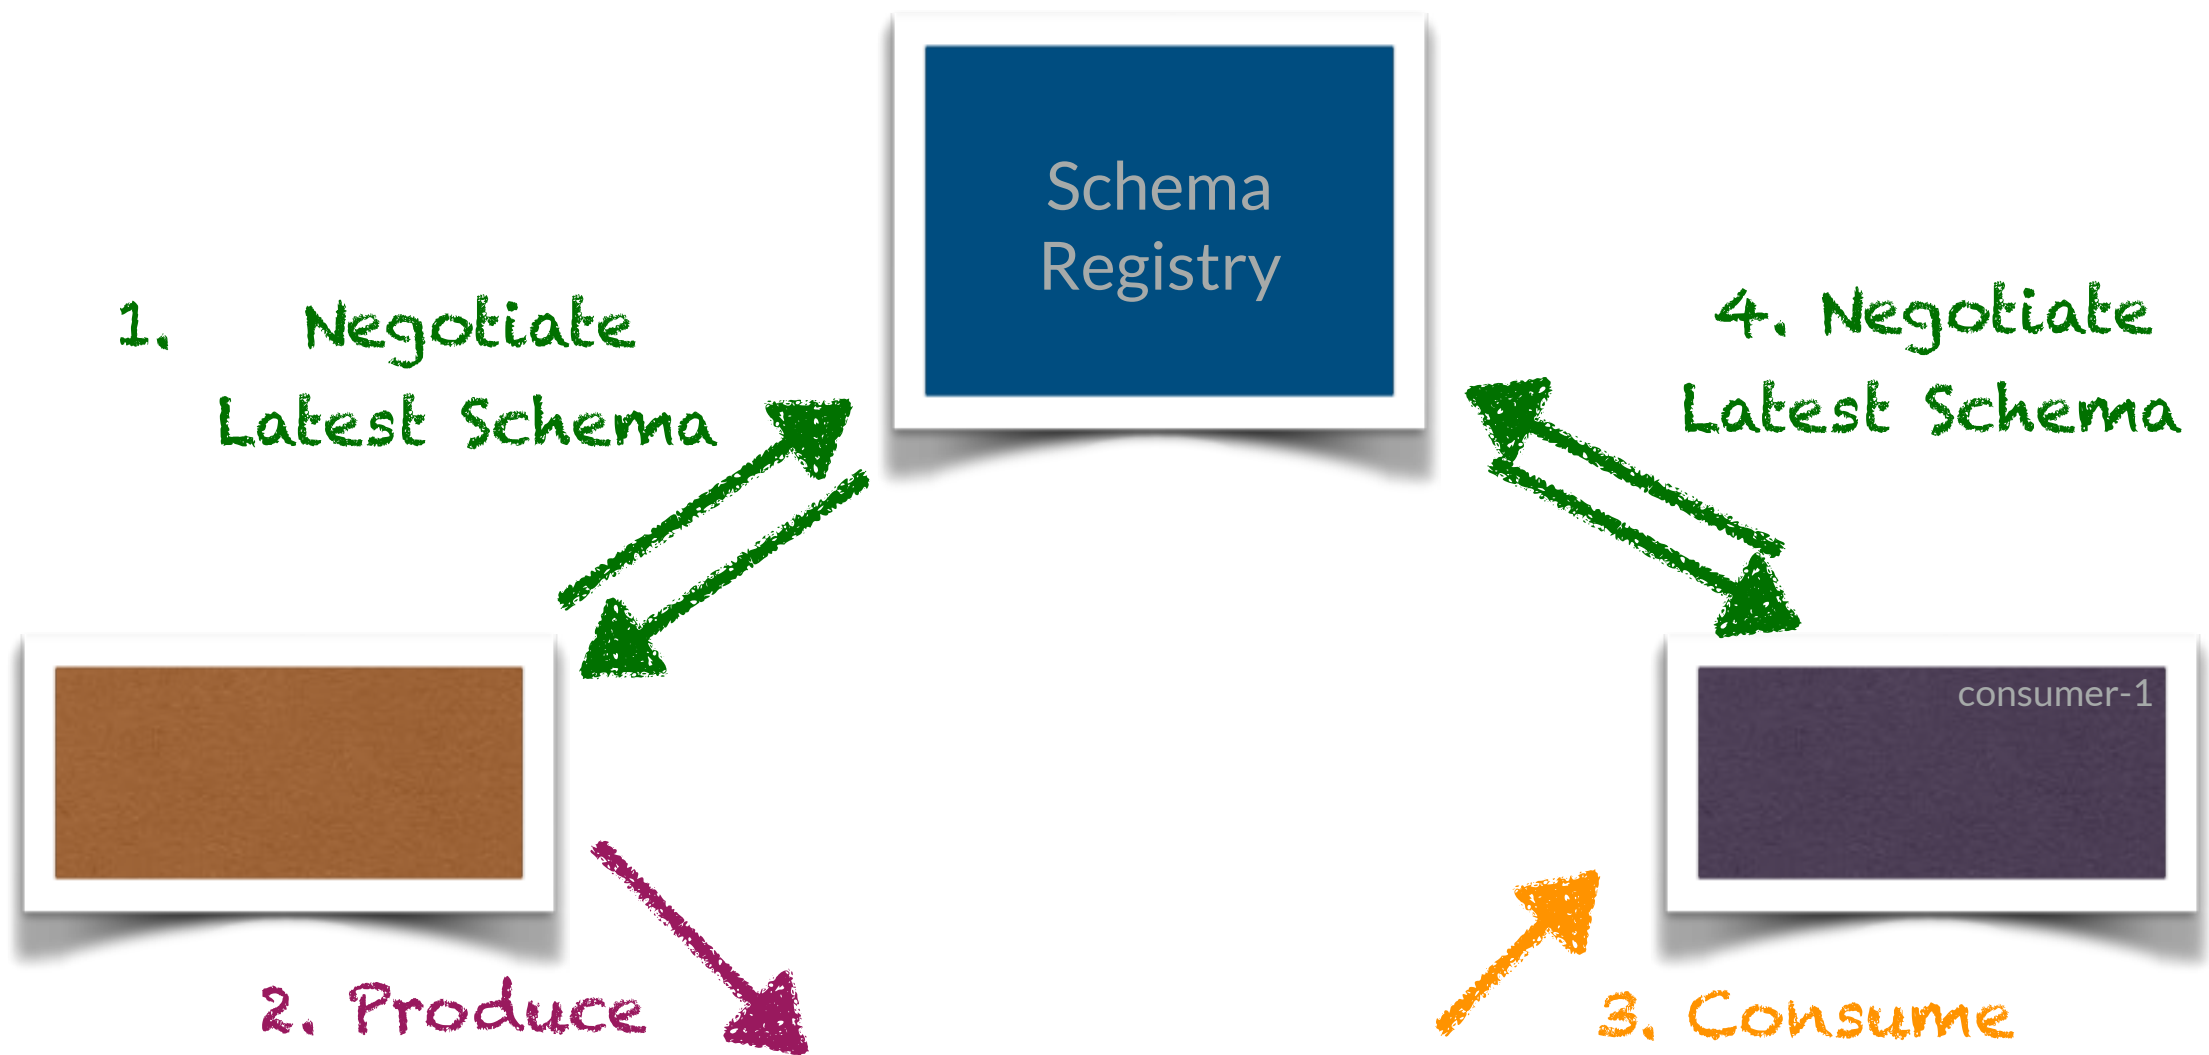




Kafka Serialization Demo



Confluent Platform





Thank You

Email: dhinojosa@evolutionext.com

Twitter: @dhinojosa